

---

# CS394R: Reinforcement Learning: Theory and Practice -- Spring 2022

Instructors: [Scott Niekum](#) and [Peter Stone](#)

Department of Computer Science

Tuesday, Thursday 9:30-11:00am

[GDC 2.216](#)

---

This course has a large capacity, so there should be room for students in other departments to register once the semester is underway.

If you're interested in taking the course, please plan on coming to the initial class sessions. We intend to post information on how to do so on piazza.

---

## IMPORTANT:

**The first reading assignment is due by class time on Tuesday 1/18. Please read chapter one of the textbook (available online)** Section 1.7 is optional. Details are available on the assignments page and below.

The second reading assignment **and written response to the reading** is due on Wednesday 1/19 at 5pm.

**Please register for the course on [edX Edge](#) (available on Monday, 1/17) and sign up for the [Piazza page](#) (if asked for an access code, use "cs394r")**

---

Jump to the [assignments page](#).

Jump to the [resources page](#).

Jump to the [textbook page](#).

Jump to the [edX course page](#).

Jump to the [project page](#).

---

## Instructor Contact Information

### Scott Niekum

**office hours:** Tuesdays 3pm-4pm and by appointment

**office:** GDC 3.404

**phone:** 232-74741

**fax:** 471-8885

**email:** [sniekum@cs.utexas.edu](mailto:sniekum@cs.utexas.edu)

### Peter Stone

**office hours:** Tuesdays 1:10pm-2pm and by appointment

**office:** GDC 3.508

**phone:** 471-9796

**fax:** 471-8885

**email:** [pstone@cs.utexas.edu](mailto:pstone@cs.utexas.edu)

---

## Teaching Assistants

### [William Macke](#)

**office hours:** Thursday 3-4 PM and Friday 1-2PM CST

**email:** [wmacke@utexas.edu](mailto:wmacke@utexas.edu)

[Zoom Link](#)

### [Harshit Sikchi](#)

**office hours:** Wednesday 3-4 PM CST (on zoom)

**office:** GDC 3.408B

**email:** [hsikchi@utexas.edu](mailto:hsikchi@utexas.edu)

[Zoom Link](#)

### **Rohan Nair**

**office hours:** Monday and Wednesday 2-3 PM CST

**email:** [rohan.nair@utexas.edu](mailto:rohan.nair@utexas.edu)

[Zoom Link](#)

### **Zizhao Wang**

**office hours:** Thursday 2-3 PM CST (on zoom)

**email:** [zizhao.wang@utexas.edu](mailto:zizhao.wang@utexas.edu)

[Zoom Link](#)

---

## Course Description

*"The idea that we learn by interacting with our environment is probably the first to occur to us when we think about the nature of learning. When an infant plays, waves its arms, or looks about, it has no explicit teacher, but it does have a direct sensori-motor connection to its environment. Exercising this connection produces a wealth of information about cause and effect, about the consequences of actions, and about what to do in order to achieve goals. Throughout our lives, such interactions are undoubtedly a major source of knowledge about our environment and ourselves. Whether we are learning to drive a car or to hold a conversation, we are all acutely aware of how our environment responds to what we do, and we seek to influence what happens through our behavior. Learning from interaction is a foundational idea underlying nearly all theories of learning and intelligence."*

*"Reinforcement learning problems involve learning what to do --- how to map situations to actions --- so as to maximize a numerical reward signal. In an essential way these are closed-loop problems because the learning system's actions influence its later inputs. Moreover, the learner is not told which actions to take, as in many forms of machine learning, but instead must discover which actions yield the most reward by trying them out. In the most interesting and challenging cases, actions may affect not only the immediate reward but also the next situation and, through that, all subsequent rewards. These three characteristics --- being closed-loop in an essential way, not having direct instructions as to what actions to take, and where the consequences of actions, including reward signals, play out over extended time periods --- are the three most important distinguishing features of the reinforcement learning problem."*

These two paragraphs from chapter 1 of the course textbook describe the topic of this course. The course is a graduate level class. There will be assigned readings and class discussions and activities. It will cover the first 13 chapters of the (2nd edition of the) course textbook plus Chapter 16. Beyond that, we will move to more

advanced and/or recent readings from the field with an aim towards focussing on the practical successes and challenges relating to reinforcement learning.

In addition to the readings, there will be one exam, some problem sets, programming assignments, and a final project. Students will be expected to be proficient programmers.

---

## Prerequisites

Strong programming skills and knowledge of probability are required. Some background in artificial intelligence is recommended.

---

## Text

The course textbook is:

[Reinforcement Learning: An Introduction](#).

By [Richard S. Sutton](#) and [Andrew G. Barto](#).

MIT Press, Cambridge, MA, 2018.

Note that the book is available [on-line](#), though if you take the course, it's probably a book you'll want for your bookshelf.

---

## Assignments

Reading, written, and programming assignments will be updated on the [assignments page](#). A tentative schedule for the entire semester is posted. But the readings and exercises may change up until the Wednesday of the week before they are due (1 week in advance).

---

## Resources

Slides from class and other relevant links and information are on the [resources page](#). If you find something that should be added there, please email it to the instructors and/or TAs.

---

## Discussion Forum

While the Professor and the TA would be glad to answer any questions you have, you would frequently find your peers to be an equally important resource in this class.

Please subscribe to our class [piazza](#) page.

---

## Course Requirements

Grades will be based on:

**Written responses to the readings and other class participation (10%):**

By 5pm on the afternoon before a class with a new reading assignment due, everyone must submit a brief question or comment about the readings in the **Readings Response** section on the [edX course page](#). Please include your name and eid in the response. In some cases, specific questions may be posted along with the readings. But in general, it is free form. Credit will be based on evidence that you have done the readings carefully. Acceptable responses include (but are not limited to):

- Insightful questions;
- Clarification questions about ambiguities;
- Comments about the relation of the reading to previous readings;
- Solutions to problems or exercises posed in the readings;
- Critiques;
- Thoughts on what you would like to learn about in more detail;
- Possible extensions or related studies;
- Thoughts on the paper's importance; and
- Summaries of the most important things you learned.

Example successful responses from a previous class are available on the [sample responses page](#).

These responses will be graded on a 10-point scale with a grade of 10 being a typical full-credit grade for a reasonable response. Responses will be due by 5pm on Monday. No late responses will be accepted.

This deadline is designed both to encourage you to do the readings before class and also to allow us to incorporate some of your responses into the class discussions.

Students are expected to be present in class having completed the readings and participate actively in the discussions and activities.

**Multiple choice and short answer exercises (15%):**

There will be a series of multiple choice and/or short answer questions on EdX to complete during the first 8 weeks of the semester. You will be able to submit these multiple times until you get the right answer, so they will essentially become a completion grade. But the midterm will be a similar format to these questions (though somewhat more difficult), so you will benefit from completing them carefully and making sure you fully understand the answers.

**Programming exercises (25%):**

Each student will be required to complete a series of minor programming assignments. These exercises will not involve extensive or elaborate programs. The emphasis is to be on empirically analyzing various learning algorithms and reporting on the results. They will be auto-graded. Details are on the class EdX page.

For the programming assignments, students may not use any example code found on the web or from any other source, especially for the concepts that are being covered by the assignment. If there is general purpose "plumbing" code that you would like to use, please check first with the course staff.

**Midterm Exam (25%):**

There will be a midterm exam during week 9 of the semester, covering the material from the textbook.

**Final Project (25%):**

There will be a [final project](#) due by **Dec TBA at 11:59pm** with an extension to **Dec TBA at 11:59pm**. For each day late, 1 point out of 100 will be deducted from the project grade.

## Extension Policy

If you turn in your assignment late, expect points to be deducted. No exceptions will be made for the written responses to readings-based questions (subject to the ``notice about missed work due to religious holy days" below). For other assignments, TBA.

The greater the advance notice of a need for an extension, the greater the likelihood of leniency.

---

## Academic Dishonesty Policy

You are encouraged to discuss the readings and concepts with classmates. But all written work must be your own. And programming assignments must be your own except for 2-person teams when teams are authorized. All work ideas, quotes, and code fragments that originate from elsewhere must be cited according to standard academic practice. Students caught cheating will automatically fail the course. If in doubt, look at the [departmental guidelines](#) and/or ask.

---

## Notice about students with disabilities

The University of Texas at Austin provides upon request appropriate academic accommodations for qualified students with disabilities. To determine if you qualify, please contact the Dean of Students at 471-6529; 471-4641 TTY. If they certify your needs, I will work with you to make appropriate arrangements.

---

## Notice about missed work due to religious holy days

A student who misses an examination, work assignment, or other project due to the observance of a religious holy day will be given an opportunity to complete the work missed within a reasonable time after the absence, provided that he or she has properly notified the instructor. It is the policy of the University of Texas at Austin that the student must notify the instructor at least fourteen days prior to the classes scheduled on dates he or she will be absent to observe a religious holy day. For religious holy days that fall within the first two weeks of the semester, the notice should be given on the first day of the semester. The student will not be penalized for these excused absences, but the instructor may appropriately respond if the student fails to complete satisfactorily the missed assignment or examination within a reasonable time after the excused absence.

---

## Related Courses

- [Previous offering of this course](#) (Fall 2019)
- [The one before that](#) (Fall 2016)
- [The one before that](#) (Spring 2013)
- [The one before that](#) (Spring 2011)
- [The one before that](#) (Fall 2007)
- [The one before that](#) (Fall 2004)

- [Emma Brunskill's course at Stanford](#)
  - [Michael Littman's course at Brown](#)
  - [Sergey Levine's course at Berkeley](#)
  - [David Silver's course at UCL](#)
  - [Yishay Mansour's course at Tel Aviv University](#)
- 

## UTCS Reinforcement Learning Reading Group

- The UTCS Reinforcement Learning Reading Group is a student run group that meets bi-weekly to discuss papers related to reinforcement learning. The [RL Reading Group web page](#) also provides a repository of past readings.
  - Here's An [RL reading list](#) from Shivaram Kalyanakrishnan.
  - Csaba Szepesvari's list of [RL applications](#)
- 

[\[Back to Department Homepage\]](#)

*Page maintained by [Peter Stone](#)  
Questions? Send me [mail](#)*